### **Computing Infrastructure**

- Connecting to Adroit using ssh
- Software environment using module
- Software versioning using git
- Compiling using make
- Compiling and running using MPI and slurm
- Plotting using gnuplot, IDL and python
- Bash scripting

### adroit.princeton.edu

Install an X server and a terminal window on your laptop.

- for Mac OS:
  - iTerm for the terminal window
  - XQuartz for the X server
- for Windows:
  - PuTTY or KiTTY to connect remotely
  - Xming for the X server.

Get an account on adroit at <a href="https://researchcomputing.princeton.edu">https://researchcomputing.princeton.edu</a>

Connect to <u>adroit.princeton.edu</u> with 2 options:

- off campus, connect first to the VPN (netID and password).
- in the terminal window, enter: ssh -X yourlogin@adroit.princeton.edu
   Enter your password (again...)
- off campus, connect first to the VPN.
- connect to <u>myadroit.princeton.edu</u> using your favorite browser.

## **Configuring ssh**

The SSH protocol uses public key cryptography for authenticating hosts and users. The authentication keys, called SSH keys, are created using the keygen program.

On your local laptop, enter ssh-keygen without any arguments (just hit return).

It will create the directory .ssh with the private key stored in file id\_rsa and the public key stored in file id\_rsa.pub

Copy / paste the public key into a new file called authorized\_keys into the remote .ssh directory. Make store the public key is stored using only one line (removing all line breaks).

You will not be asked for your password anymore (using ssh and scp).

Enter ssh -X netid@adroit.princeton.edu to allow X11 forwarding to run graphical applications remotely.

You can also edit your SSH config file. First enter the command cp /etc/ssh/ssh\_config ~/.ssh/config In this new file, set the parameter X11Forwarding yes

### Set up your adroit environment

Typically users initialize their environment when they log in by setting environment information for every application they will reference during the session.

The Environment Modules package is a tool that simplify shell initialization and lets users easily modify their environment during the session with modulefiles.

> module avail

[rt3504@adroit5 ~]\$ module avail dot module-git module-info modules null use.own boost/1.73.0 cudatoolkit/11.4 qs1/2.6netcdf/gcc/hdf5-1.10.6/4.7.3 netcdf/gcc/hdf5-1.10.6/openmpi-4.1.0/4.7.4 netcdf/gcc/hdf5-1.10.6/4.7.4 cmake/3.18.2fftw/gcc/3.3.9 hdf5/gcc/1.10.6 openmpi/gcc/4.1.0 cudatoolkit/10.2 fftw/gcc/intel-mpi/3.3.9 hdf5/gcc/intel-mpi/1.10.6 netcdf/gcc/hdf5-1.10.6/intel-mpi/4.7.3 R/3.6.3cudatoolkit/11.1 fftw/gcc/openmpi-4.1.0/2.1.5 hdf5/gcc/openmpi-4.0.4/1.10.6 netcdf/gcc/hdf5-1.10.6/intel-mpi/4.7.4 R/4.0.5cudatoolkit/11.3 fftw/gcc/openmpi-4.1.0/3.3.9 hdf5/gcc/openmpi-4.1.0/1.10.6 netcdf/gcc/hdf5-1.10.6/openmpi-4.0.4/4.7.3 ucx/1.9.0------/opt/share/Modules/modulefiles -course/cee306/default intel-dpl/2021.1.2 intel-mpi/qcc/2019.7 intel-rt/2021.1.2 course/pyapu/default intel-vtune/oneapi intel/2021.1.2 intel-inspector/oneapi intel-mpi/gcc/2021.3.1 course/elecos396/default course/pyJC/default intel-tac/2020.1-024 intel/19.1(@) java/11 course/env-geo-367/default intel-advisor/oneapi intel-mpi/intel/2019.7 intel-mkl/2020.1 intel-tac/2021.1.1 intel/19.1.1.217 nvhpc-nocompiler/21.5 course/pycompiled/default intel-debugger/10.0.0 intel-mkl/2021.1.1 intel-mpi/intel/2021.3.1 intel-tbb/2021.1.1 intel/2021.1(@) nvhpc/21.5-----/usr/licensed/Modules/modulefiles ansys/AnsysEM21.1 gurobi/9.0.1 julia/0.7.0 lumerical/2020R2 matlab/R2011a matlab/R2016a matlab/R2020b spark/hadoop2.7/2.2.0 stata/16.1(defail agilent/2017 agilent/latest cadence/20171206 gurobi/9.1.2 julia/1.0.1 map/18.0.2 matlab/R2011b matlab/R2016b matlab/R2021a spark/hadoop2.7/2.4.6 stata/17.0 anaconda3/2018.12 cadence/latest IDL/8.4 spark/hadoop3.2/3.2.0 synopsys/saed-m julia/1.0.3 map/20.0.1 matlab/R2012a matlab/R2017a paraview/5.9.1 anaconda3/2019.3 ddt/18.0.2IDL/8.5 julia/1.1.0 mathematica/11.0 matlab/R2013a matlab/R2017b(default) performance-reports/7.1 stata/11.0 anaconda3/2019.10 ddt/20.0.1IDL/8.6 julia/1.2.0 mathematica/11.1.0 matlab/R2013b matlab/R2018a spark/hadoop1/1.6.1 stata/12.0 anaconda3/2020.2 julia/0.4.7 julia/1.3.0 mathematica/11.3.0 matlab/R2014a matlab/R2018b spark/hadoop2.6/1.5.2 stata/13.0 EDEM/2020.0 julia/0.5.0 julia/1.4.1 mathematica/12.1.1 matlab/R2014b matlab/R2019a spark/hadoop2.6/1.6.1 anaconda3/2020.7 gurobi/7.5.1 stata/14.0 julia/1.5.0 matlab/R2010a spark/hadoop2.6/2.0.0 anaconda3/2020.11 qurobi/8.0.1(default) julia/0.5.1 matlab/R2015a matlab/R2019b stata/15.0 anaconda3/2021.5 gurobi/8.1.1 julia/0.6.0 julia/1.6.1 matlab/R2010b matlab/R2015b matlab/R2020a spark/hadoop2.6/2.1.0 stata/16.0

## Set up your adroit environment

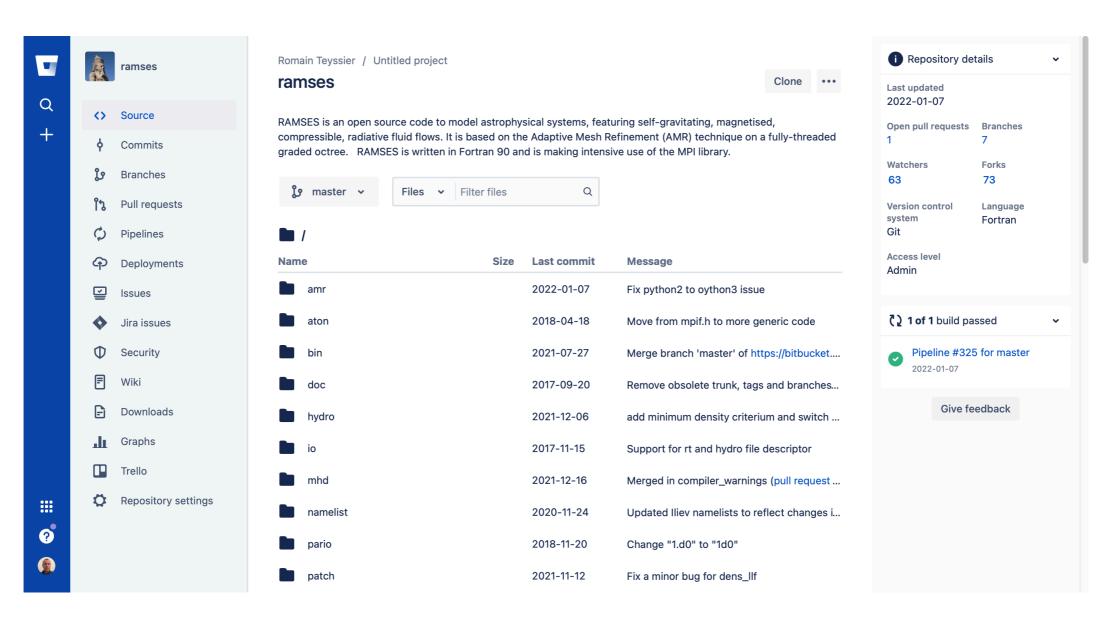
```
> module load, module list
[rt3504@adroit5 ~]$ module list
No Modulefiles Currently Loaded.
[rt3504@adroit5 ~]$ module load openmpi/gcc/4.1.0
[rt3504@adroit5 ~]$ module load fftw/gcc/openmpi-4.1.0/3.3.9
[rt3504@adroit5 ~]$ module load anaconda3/2021.5
(base) [rt3504@adroit5 ~]$ module list
Currently Loaded Modulefiles:
1) openmpi/gcc/4.1.0 2) fftw/gcc/openmpi-4.1.0/3.3.9 3) anaconda3/2021.5
> module switch
[rt3504@adroit5 ~]$ ifort
-bash: ifort: command not found
[rt3504@adroit5 ~]$ module load intel/2021.1
Loading intel/2021.1.2
  Loading requirement: intel-rt/2021.1.2 intel-tbb/2021.1.1 intel-mkl/2021.1.1
    intel-debugger/10.0.0 intel-dpl/2021.1.2
    /opt/intel/oneapi/compiler/2021.1.2/linux/lib/oclfpga/modulefiles/oclfpga
[rt3504@adroit5 ~]$ ifort
ifort: command line error: no files specified; for help type "ifort -help"
[rt3504@adroit5 ~]$ which ifort
/opt/intel/oneapi/compiler/2021.1.2/linux/bin/intel64/ifort
[rt3504@adroit5 ~]$ module switch intel/2021.1.2 intel/19.1
Switching from intel/2021.1.2 to intel/19.1.1.217
  Unloading useless requirement: intel-mkl/2021.1.1 intel-rt/2021.1.2
     intel-tbb/2021.1.1 intel-debugger/10.0.0 intel-dpl/2021.1.2
    /opt/intel/oneapi/compiler/2021.1.2/linux/lib/oclfpga/modulefiles/oclfpga
  Loading requirement: intel-mkl/2020.1
[rt3504@adroit5 ~]$ which ifort
/opt/intel/compilers_and_libraries_2020.1.217/linux/bin/intel64/ifort
```

## Source Code Management with Git

Create your account on GitHub or BitBucket.

Git is a version control system that developers use all over the world. It helps you track different versions of your code and collaborate with other developers.

In your browser, enter: <a href="https://bitbucket.org/rteyssie/ramses">https://bitbucket.org/rteyssie/ramses</a>



## Source Code Management with Git

Clone the code to your account on adroit.

For this, in a terminal window, enter:

> git clone https://rteyssie@bitbucket.org/rteyssie/ramses.git

You have now the latest version of the RAMSES code in your HOME directory.

You can update later changes using git pull.

If you have write authorization, you can also upload changes using git push.

git add foo.c Add a file in the local git repository.

git rm foo.c Remove a file from the local git repository.

git commit -a -m "text describing the changes"

Commit changes to the local git repository.

Upload all local changes to the remote repository.

git push

## **Source Code Management with Git**

What you can do with a Git version control system.

- track bugs to their authors
- edit users guides and wikis
- revert to older versions
- track the exact version of the code with the Git hash key
- set up an automatic build and test system
- validation and verification (reproducibility)

See examples in <a href="https://bitbucket.org/rteyssie/ramses">https://bitbucket.org/rteyssie/ramses</a>

### **Compiling Code with Makefile**

Makefiles are a simple way to automate code compilation.

A simple makefile contains a set of "rules" with the following syntax:

target: dependencies

command1 command2

A target is the name of a file that is generated by a program such as executable or object files. A target can also be the name of an action to carry out, such as `clean'.

A *dependency* is a file that is used as input to create the target. A target often depends on several files. Only files that have been recently modified are recompiled.

A *command* is an action that make carries out. A rule may have more than one command, each on its own line. You need to put a tab character at the beginning of every command line.

#### A complex example here:

> cd ramses/bin

> make

> make clean

#### A simple example here:

> cd ramses/utils/f90

> make

> make clean

## **Compiling Code with Makefile**

```
[rt3504@adroit5 f90]$ more Makefile
# Makefile for RAMSES utils
BINDIR=.
# gfortran configuration
F90=afortran
CFLAGS=-w -ffree-line-length-none -std=f2008 -Ofast -Wall -x f95-cpp-input
                                                                                     local variables
# ifort configuration
#F90=ifort
#CFLAGS=-cpp
LFLAGS=
MODOBJ=random.o utils.o io_ramses.o
# Make all targets
all: amr2prof amr2cylprof ramses2tipsy amr2map part2map part2prof \
                                                                                  make all or make
     part2cube part2list partcenter part2birth part2sfr getstarlist \
     part2cylprof
amr2prof: $(MODOBJ) amr2prof.o
       $(F90) $(LFLAGS) $^ -o $(BINDIR)/$@
amr2cylprof: $(MODOBJ) amr2cylprof.o
                                                           targets, dependencies and actions
        $(F90) $(LFLAGS) $^ -o $(BINDIR)/$@
amr2map: $(MODOBJ) amr2map.o
        $(F90) $(LFLAGS) $\(^-\)o $(BINDIR)/$@
# Make a specific object file
%.o: %.f90
       $(F90) $(CFLAGS) -c $^ -o $@
                                                      wild cards
clean:
       rm *.o *.mod
```

phony action (clean)

# Running the code and Visualize with gnuplot

Once the code is compiled, run it with a parameter file as command

line argument.

```
> bin/ramses1d namelist/tube1d.nml > run.log
```

The ket symbol > redirects the standard output to a file. This simple example runs a standard hydro test case, the Sod shock tube.

You can visualize quickly the results using gnuplot. Enter: > gnuplot

You are now using the gnuplot plotting program.

Make sure you have properly setup X11 forwarding and enter:

gnuplot> plot "run.log" every ::::1000 u 2:3 title "initial density"

gnuplot> replot "run.log" every ::1000 u 2:3 title "final density"

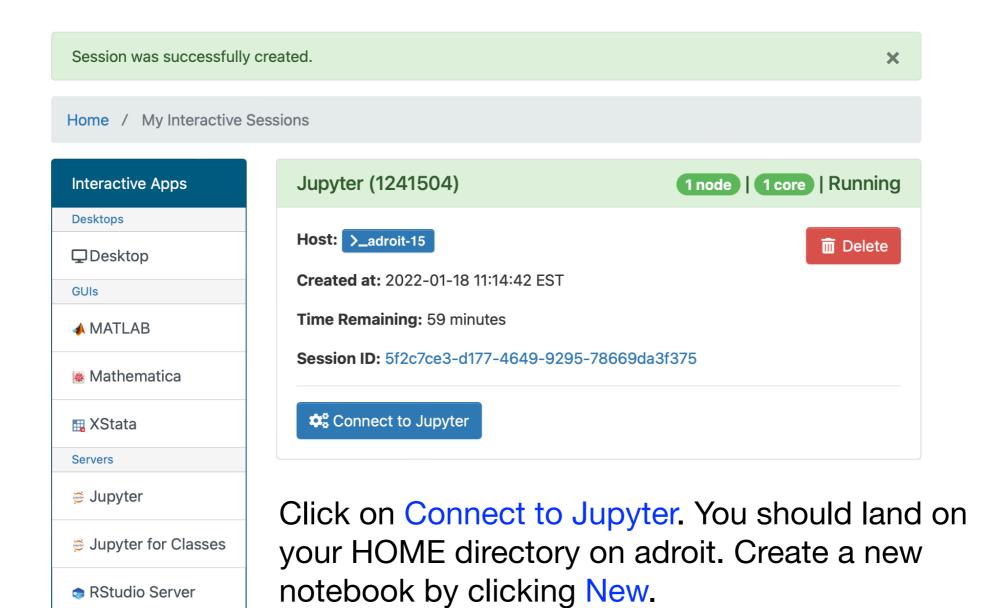
```
You can also use gnuplot without X11 on the command line with:
```

```
> gnuplot -e "set term jpeg; plot 'run.log' every :: 1000 u 2:3 title 'final density' " > density.jpeg
```

### Visualize with a Jupyter Notebook

On your browser (after launching the VPN) connect to the web site: <a href="https://myadroit.princeton.edu/">https://myadroit.princeton.edu/</a>

Launch a Jupyter Server interactive session (click Launch). After a few seconds, you should see this:



### Visualize using jupyter notebook

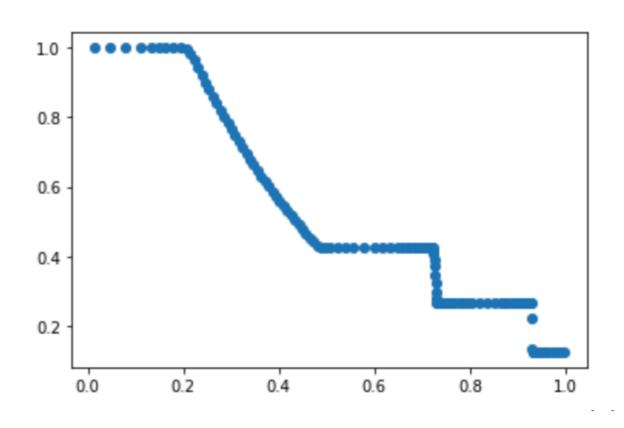
Edit your own notebook following this example.

If you are missing a package, you need to add it on adroit. Enter (for example) the following:

```
> module load anaconda3/2021.5
> pip3 install astropy
```

Ultimately, you should see this plot:

More details and options can be found at <a href="https://researchcomputing.princeton.edu/support/knowledge-base/jupyter">https://researchcomputing.princeton.edu/support/knowledge-base/jupyter</a>



## **Compiling and Running with MPI**

The Message Passing Library allows you to run programs in parallel on multiple processors.

```
> module load openmpi/gcc/4.1.0
> cd ramses/bin
> make clean
> make NDIM=2 MPI=1
```

#### Launch a parallel execution using slurm.

```
> srun -n 8 -t 00:01:00 bin/ramses2d namelist/sedov2d.nml > run.log
```

### You can also edit a slurm script in a file called (for example) job.sh

```
#!/bin/bash -l
#SBATCH --job-name=sedov2d
#SBATCH --time=00:01:00
#SBATCH --nodes=1
#SBATCH --ntasks=8

module load openmpi/gcc/4.1.0
export DATE=`date +%F_%Hh%M`
srun ~/ramses/bin/ramses2d ~/ramses/namelist/sedov2d.nml > run_$DATE.log
```

## **Compiling and Running with MPI**

#### Slurm allows you to submit a job on the queue

```
> sbatch job.sh
```

#### You can check your status on the queue

```
> squeue -u yourlogin

JOBID PARTITION NAME USER ST TIME NODES NODELIST(REASON)

1241531 all sedov2d rt3504 PD 0:00 1 (Resources)

1241532 all sedov2d rt3504 R 0:01 1 adroit-08
```

#### You can cancel your job on the queue.

> scancel 1241531

#### You can find more information here:

https://researchcomputing.princeton.edu/support/knowledge-base/slurm

Please use /scratch/network/yourlogin for the output of running jobs and to store large datasets. This space is an NFS-mounted shared space of close to 24 TB. Run the checkquota command to see your usage.

### **BASH** scripting

First, run a new parallel 2D simulation with much more output files.

```
> cp namelist/sedov2d.nml sedov2d.nml
```

#### Modify this new file in the namelist block

```
&OUTPUT_PARAMS
tend=1.0
delta_tout=0.01
```

#### Launch a parallel execution using slurm.

```
> srun -n 8 -t 00:01:00 bin/ramses2d sedov2d.nml
```

### You should see now 101 output directories.

Now edit a new bash script: > emacs -nw movie.sh

```
#!/bin/bash
for i in `ls -d output_00*`
do
    echo $i
    ./utils/f90/amr2map -inp $i -out dens.map -typ 1 -fil ascii -lma 8
    gnuplot -e "set term jpeg; set palette rgbformulae 22,13,-31; set pm3d map; set size square;
set cbrange [0:7]; splot 'dens.map' u 1:2:3 notitle" > pic_$i.jpeg
done
echo "converting to animated gif"
convert -delay 1 pic_output_00* movie.gif
rm -rf pic_output_00*
```

Make sure the script is executable: > chmod a+x movie.sh

Execute by entering: > ./movie.sh

### **BASH** scripting

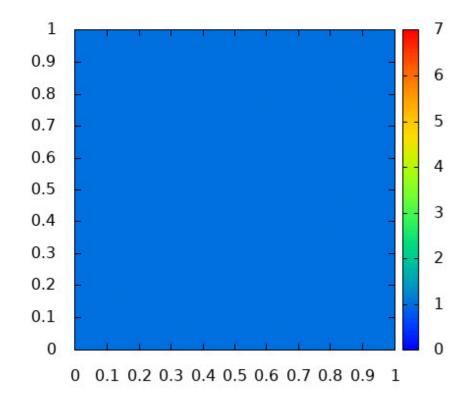
#### A variant for the main loop:

```
#!/bin/bash
for i in `seq -w 101`  # be careful these are backquotes or backticks
do
        echo $i
        ~/ramses/utils/f90/amr2map -inp output_00$i -out dens.map -typ 1 -fil ascii -lma 8
        gnuplot -e "set term jpeg; set palette rgbformulae 22,13,-31; set pm3d map; set size square;
set cbrange [0:7]; splot 'dens.map' u 1:2:3 notitle" > pic_00$i.jpeg
done
echo "converting to animated gif"
convert -delay 1 pic_00* movie.gif
rm -rf pic_00*
```

#### Copy the animated gif back to your laptop

> scp yourlogin@adroit.princeton.edu:ramses/movie.gif .

#### Read it with your favorite web browser. It should look like that.



#### Other useful Linux commands:

- echo: print to standard output
- grep: find a string in a file line by line
- cut: extract a string from a line
- sort: sort column according to some rules
- sed: advanced fine edition
- awk: advanced math operations
- man: manual of Linux command